

Decentralized Crash-Resilient Runtime Verification

Shokoufeh Kazemlou

Department of Computing and Software
McMaster University

December 13, 2017

Overview

- 1 Motivation
- 2 LTL₃ Monitor
- 3 Problem Statement
- 4 Related Work
- 5 Contributions
- 6 Synchronous Monitoring
- 7 Asynchronous Monitoring
- 8 Conclusion
- 9 Future Work

Presentation outline

- 1 Motivation
- 2 LTL₃ Monitor
- 3 Problem Statement
- 4 Related Work
- 5 Contributions
- 6 Synchronous Monitoring
- 7 Asynchronous Monitoring
- 8 Conclusion
- 9 Future Work

Motivation

Traditional Verification

Exhaustive verification methods are extremely valuable to ensure system-wide correctness.

They often require developing an abstract model of the system and may suffer from the infamous **state-explosion** problem.

Motivation

Traditional Verification

Exhaustive verification methods are extremely valuable to ensure system-wide correctness.

They often require developing an abstract model of the system and may suffer from the infamous **state-explosion** problem.

Runtime Verification

Runtime verification (RV) refers to a technique, where a monitor checks at run time whether or not the execution of a system under inspection satisfies a given correctness property.

RV **complements** exhaustive verification techniques as well as underapproximated methods such as testing and tracing.

Motivation

RV in Distributed Systems

Designing a **decentralized runtime monitor** for a **distributed system** is an especially difficult task since it deals with

- computing **global snapshots** at run time, and
- estimating the **total order** of events

in order for the monitor to reason about the temporal behavior of the system.

Presentation outline

- 1 Motivation
- 2 **LTL₃ Monitor**
- 3 Problem Statement
- 4 Related Work
- 5 Contributions
- 6 Synchronous Monitoring
- 7 Asynchronous Monitoring
- 8 Conclusion
- 9 Future Work

3-Valued LTL (LTL₃) [Bauer, Leucker, Schallhart 11]

3-valued LTL evaluates LTL formulas for finite words with an eye on **possible future extensions**.

Three Truth Values

The set of truth values is $\mathbb{B}_3 = \{\top, \perp, ?\}$, where

- **\top** : the formula is **permanently satisfied** no matter how the current execution extends,
- **\perp** : the formula is **permanently violated** no matter how the current execution extends
- **$?$** : denotes an unknown verdict; i.e., there exist extensions that can falsify or make true the formula.

3-Valued LTL

LTL₃ Semantics

Let $u \in \Sigma^*$ be a finite word. The truth value of an LTL₃ formula φ with respect to u , denoted by $[u \models_3 \varphi]$, is defined as follows:

$$[u \models_3 \varphi] = \begin{cases} \top & \text{if } \forall w \in \Sigma^\omega : uw \models \varphi \\ \perp & \text{if } \forall w \in \Sigma^\omega : uw \not\models \varphi \\ ? & \text{otherwise.} \end{cases}$$

3-Valued LTL

LTL₃ Monitor


Let φ be an LTL formula. The **LTL₃ monitor** of φ is the unique deterministic finite state machine $\mathcal{M}_3^\varphi = (\Sigma, Q, q_0, \delta, \lambda)$, where Q is a set of states, q_0 is the initial state, $\delta \subseteq Q \times \Sigma \times Q$ is the transition relation, and $\lambda : Q \rightarrow \mathbb{B}_3$, is a function such that:

$$\lambda(\delta(q_0, u)) = [u \models_3 \varphi]$$

for every finite word $u \in \Sigma^*$. □

Example

LTL₃ monitor for $a \mathbf{U} b$




Presentation outline

- 1 Motivation
- 2 LTL₃ Monitor
- 3 Problem Statement**
- 4 Related Work
- 5 Contributions
- 6 Synchronous Monitoring
- 7 Asynchronous Monitoring
- 8 Conclusion
- 9 Future Work

Problem Statement

Distributed Monitors

Let $\mathcal{M} = \{M_1, M_2, \dots, M_n\}$ be a set of **distributed monitors** monitoring an underlying system.



Synchronous Monitoring

Local Monitor Algorithm

Data: LTL formula φ and state s_j

Result: a verdict from \mathbb{B}_3

- 1 Let $S_i^{s_j}$ be the initial concrete local state of the monitor
- 2 $LS_i^1 \leftarrow \mu(S_i^{s_j}, \varphi)$
- 3 **for** $r = 1, 2, \dots$ **do**
 - 4 **Send:** broadcasts its current abstract local state LS_i^r
 - 5 **Receive:** let $\Pi_i^r = \{LS_j^r\}_{j \in [1, n]}$ be the set of all messages received at round r .
 - 6 **Computation:** $LS_i^{r+1} \leftarrow LC(\Pi_i^r)$
- 7 emits a verdict from \mathbb{B}_3

Problem Statement

A non-faulty monitor should compute and emit a verdict that a centralized monitor that has global view of the system would compute. Formally:

$$\forall i \in [1, n] : M_i \text{ is non-faulty} \rightarrow \nu_i = [\alpha \models_3 \varphi]$$

Presentation outline

- 1 Motivation
- 2 LTL₃ Monitor
- 3 Problem Statement
- 4 Related Work**
- 5 Contributions
- 6 Synchronous Monitoring
- 7 Asynchronous Monitoring
- 8 Conclusion
- 9 Future Work

Related Work

Central Monitor

- H. Chauhan and V. K. Garg and A. Natarajan and N. Mittal. **A Distributed Abstraction Algorithm for Online Predicate Detection** (SRDS 2013).
- N Mittal and V. K. Garg. **Techniques and applications of computation slicing** (Distributed Computing 2005).
- V. A. Ogale and V. K. Garg. **Detecting Temporal Logic Predicates on Distributed Computations** (DISC 2007).

Related Work

Fault-free Setting

- A. Bauer and Y. Falcone. **Decentralised LTL monitoring** (FMDS 2016).
- C. Colombo and Y. Falcone. **Organising LTL monitors over distributed systems with a global clock** (FMDS 2016).
- M. Mostafa, B. Bonakdarpour. **Decentralized Runtime Verification of LTL Specifications in Distributed Systems**. (IPDPS 2015).

Fault-tolerant Distributed Monitoring

B. Bonakdarpour, P. Fraigniaud, S. Rajsbaum, D. A. Rosenblueth, C. Travers. **Decentralized Asynchronous Crash-Resilient Runtime Verification** (CONCUR 2016).

Presentation outline

- 1 Motivation
- 2 LTL₃ Monitor
- 3 Problem Statement
- 4 Related Work
- 5 Contributions**
- 6 Synchronous Monitoring
- 7 Asynchronous Monitoring
- 8 Conclusion
- 9 Future Work

Contributions

Contributions

- An **automata-based** distributed LTL monitoring algorithm for the decentralized crash-resilient synchronous monitoring.
- Reducing the **message size** overhead from $|AP|$ per message, to $\log(m_q)$, where m_q is the number of outgoing transitions from the current monitor state in each local monitor's automaton.
- Introducing an **Extended LTL₃** monitor for synchronous/asynchronous crash-resilient monitoring.

Presentation outline

- 1 Motivation
- 2 LTL₃ Monitor
- 3 Problem Statement
- 4 Related Work
- 5 Contributions
- 6 Synchronous Monitoring**
- 7 Asynchronous Monitoring
- 8 Conclusion
- 9 Future Work



Synchronous Monitoring

Distributed Synchronous Setting

Finite trace $\alpha = s_0 s_1 \cdots s_k$;

Set of synchronous monitors $\mathcal{M} = \{M_1, M_2, \dots, M_n\}$;

Correctness property expressed by an LTL formula φ .

Algorithm Sketch

- 1 takes a **sample** from state s_j ;
- 2 **broadcasts** a message containing its current observation, and **receives** messages from other monitors;
- 3 performs a **local computation** and updates its current observation;
- 4 **emits** a truth value from \mathbb{B}_3 .

Synchronous Monitoring

Uniform Consensus

Each process proposes a value, and the processes have to collectively agree on the same value.

- **Validity:** A decided value is a proposed value.
- **Agreement:** No two processes decide different values.
- **Termination:** Every correct process decides.

Validity Specification in Synchronous Monitoring

The decided value must be the same value that a centralized monitor with full view of the system would compute.

Number of Rounds

The lower bound on the number of rounds required to consistently monitor the system is $f + 1$, where f is the total number of crashes the system can tolerate.



Synchronous Monitoring

Local Monitor Algorithm

Data: LTL formula φ and state s_j

Result: a verdict from \mathbb{B}_3

- 1 Let $S_i^{s_j}$ be the initial concrete local state of the monitor
- 2 $LS_i^1 \leftarrow \mu(S_i^{s_j}, \varphi)$
- 3 **for** $r = 1, 2, \dots$ **do**
 - 4 *Send:* broadcasts its current abstract local state LS_i^r
 - 5 *Receive:* let $\Pi_i^r = \{LS_j^r\}_{j \in [1, n]}$ be the set of all messages received at round r .
 - 6 *Computation:* $LS_i^{r+1} \leftarrow LC(\Pi_i^r)$
- 7 emits a verdict from \mathbb{B}_3

Challenges in Synchronous Monitoring

Example

Let $\varphi = \mathbf{F}(a \wedge b)$, $AP = \{a, b\}$, and $\mathcal{M} = \{M_1, M_2, M_3, M_4\}$. Suppose $s = \{a, b\}$ is the current global state of the system, and the initial samples of the monitors are as follows:

sample

	<i>a</i>	<i>b</i>
M_1	<i>true</i>	\perp
M_2	\perp	<i>true</i>
M_3	\perp	<i>true</i>
M_4	\perp	<i>true</i>

round 1

	<i>a</i>	<i>b</i>
M_1	crashed	crashed
M_2	<i>true</i>	<i>true</i>
M_3	\perp	<i>true</i>
M_4	\perp	<i>true</i>

round 2

	<i>a</i>	<i>b</i>
M_1	crashed	crashed
M_2	crashed	crashed
M_3	<i>true</i>	<i>true</i>
M_4	\perp	<i>true</i>

round 3

	<i>a</i>	<i>b</i>
M_1	crashed	crashed
M_2	crashed	crashed
M_3	<i>true</i>	<i>true</i>
M_4	<i>true</i>	<i>true</i>

Synchronous Monitoring

Decentralized
Crash-
Resilient
Runtime
Verification

Shokoufeh
Kazemlou

Motivation

LTL₃
Monitor

Problem
Statement

Related
Work

Contributions

Synchronous
Monitoring

Asynchronous
Monitoring

Conclusion

Future
Work

Challenge

If each monitor broadcasts its sample \Rightarrow the message size is $|AP|$

Reducing the Message Size Overhead

We introduce an algorithm which decreases the message size from $|AP|$ to $\log(m_q)$ where m_q is the number of outgoing transitions from monitor state q .

Automata-based Synchronous Monitoring

The General Idea

- Each local monitor M_i evaluates the input formula and computes a **possible** set of verdicts;

$$V_i = \{\delta(q, s') \mid s' \in E(S_i^s)\}$$

- At each round, each monitor M_i broadcasts its verdict set V_i , and computes a new verdict set by applying the **intersection** function on the verdict sets received from other monitors;

$$V_i^{r+1} = LC(\Pi_i^r) = \bigcap_{j \in [1, n]} \{V_j^r\} = \bigcap_{j \in [1, n]} \{V_j^r\}$$

- After $f + 1$ rounds of communication, each monitor emits the verdict that a centralized monitor that has the global view of the system would compute.



Automata-based Synchronous Monitoring

Local Monitor Algorithm

Data: LTL₃ monitor \mathcal{M}^φ and state s_j

Result: a verdict from \mathbb{B}_3

- 1 Let $S_i^{s_j}$ be the initial concrete local state of the monitor
- 2 $LS_i^1 \leftarrow \mu_2(\mu_1(S_i^{s_j}, \mathcal{M}_\varphi)) = V_i^1$
- 3 **for** $r = 1, \dots, f + 1$ **do**
 - 4 **Send:** broadcasts its current abstract local state $LS_i^r = V_i^r$
 - 5 **Receive:** let $\Pi_i^r = \{V_j^r\}_{j \in [1, n]}$ be the set of all messages received at round r .
 - 6 **Computation:** $LS_i^{r+1} \leftarrow LC_i(\Pi_i^r) = \bigcap_{j \in [1, n]} \{V_j^r\}$
- 7 emit $\lambda_e(v_i)$

Automata-based Synchronous Monitoring

Example

Let $\varphi = \mathbf{F}(a \wedge b)$.




Figure: LTL₃ monitor of $\varphi = \mathbf{F}(a \wedge b)$.

And suppose the initial samples of the monitors are as follows:

	sample		
	a	b	V_i^1
M_1	<i>true</i>	\perp	$\{q_0, q_T\}$
M_2	\perp	<i>true</i>	$\{q_0, q_T\}$
M_3	\perp	\perp	$\{q_0, q_T\}$
M_4	\perp	\perp	$\{q_0, q_T\}$

Automata-based Synchronous Monitoring

Example

round 1		round 2		round 3	
	V_i^2		V_i^3		V_i^4
M_1	crashed	M_1	crashed	M_1	crashed
M_2	$\{q_0, q_\top\}$	M_2	crashed	M_2	crashed
M_3	$\{q_0, q_\top\}$	M_3	$\{q_0, q_\top\}$	M_3	$\{q_0, q_\top\}$
M_4	$\{q_0, q_\top\}$	M_4	$\{q_0, q_\top\}$	M_4	$\{q_0, q_\top\}$

As we see $|V_3| = |V_4| > 1$. Therefore, they cannot emit the correct verdict \top as we have $[\{a, b\} \models_3 \mathbf{F}(a \wedge b)] = \top$.

Insufficiency of LTL₃ Monitor

The LTL₃ monitor of $\varphi = \mathbf{F}(a \wedge b)$ is not sufficient to distinguish the correct verdict when local monitors have partial view of the system.

Automata-based Synchronous Monitoring

Extended LTL₃ Monitor

Input: LTL₃ monitor $\mathcal{M}^\varphi = \{\Sigma, Q, q_0, \delta, \lambda\}$

output: Extended LTL₃ monitor $\mathcal{M}_e^\varphi = \{\Sigma, Q_e, q_0, \delta_e, \lambda_e\}$

Where,

$$Q \subseteq Q_e, q_0$$

q_0 is the initial state

$\delta_e : Q_e \times \Sigma \rightarrow 2^{Q_e}$ is a transition function

$\lambda_e : Q_e \rightarrow \mathbb{B}_3$ is a mapping function, such that:

- 1 for every non-empty finite trace $\alpha \in \Sigma^*$, we have
 $\lambda_e(\delta_e(q_0, \alpha)) = \lambda(\delta(q_0, \alpha))$.
- 2 at every $q \in Q_e$ we have $|\mathcal{I}^q| = 1$.

Extended LTL₃ Monitor Construction

Transition

A transition t_j^i from monitor state q_i to monitor state q_j is defined as follows:

$$t_j^i = \{s \in \Sigma \mid \delta(q_i, s) = q_j\}$$

Indistinguishable Transitions

We say a transition t_1 is **indistinguishable** from another transition t_2 , and denote it by $indisting?(t_1, t_2)$, if the following holds:

$$\exists s \in t_2. covered?(s, t_1)$$

Covered State

We say state s is **covered** by transition t , and we denote it by $covered?(s, t)$, if we have:

$$\forall ap \in AP. \exists s' \in t. (ap \in s \Leftrightarrow ap \in s')$$

Extended LTL₃ Monitor Construction

```

Input:  $\mathcal{M}^P = \{\Sigma, Q, q_0, \delta, \lambda\}$ 
Output:  $\mathcal{M}_e^P = \{\Sigma, Q_e, q_0, \delta_e, \lambda_e\}$ 

1  $Q_e \leftarrow Q$ 
2 for every  $q_i \in Q$  do
3   Obtain the set of outgoing transitions  $T_i$  from monitor state  $q_i$ 
4   for every  $t_j^i \in T_i$  do
      /*  $t_j^i = \{s \in \Sigma \mid \delta(q_i, s) = q_j\}$  */
      /*  $N_j$  denotes the number of transitions from which  $t_j^i$  is indistinguishable, and  $K_j$  denotes the number of transitions
         indistinguishable from  $t_j^i$  */
       $N_j \leftarrow 0$ ,  $K_j \leftarrow 0$ 
      for every  $t_k^i \in T_i \setminus \{t_j^i\}$  do
        if indisting?( $t_j^i, t_k^i$ ) then
           $N_j \leftarrow N_j + 1$ 
        if indisting?( $t_k^i, t_j^i$ ) then
           $K_j \leftarrow K_j + 1$ 
      if  $N_j > 0$  then
         $\{t_{j1}^i, t_{j2}^i\} \leftarrow SPLIT(t_j^i, N_j, K_j, T_i)$ 
         $T_i \leftarrow \{t_{j1}^i, t_{j2}^i\} \cup T_i \setminus \{t_j^i\}$ 
         $Q_e \leftarrow \{q_{j1}, q_{j2}\} \cup (Q_e \setminus \{q_j\})$ 
        if  $i \neq j$  then
          for every  $t_k^i \in T_i$  do
             $\delta(q_i, s) = q_k$  for every  $s \in t_k^i$ 
             $\delta(q_{j1}, s) \leftarrow \delta(q_i, s)$  for every  $s \in \Sigma$ 
             $\delta(q_{j2}, s) \leftarrow \delta(q_i, s)$  for every  $s \in \Sigma$ 
          if  $i = j$  then
            for every  $t_k^i \in T_i$  do
               $\delta(q_{j1}, s) = q_k$  for every  $s \in t_k^i$ 
               $\delta(q_{j2}, s) \leftarrow \delta(q_{j1}, s)$  for every  $s \in \Sigma$ 
             $\lambda_e(q_{j1}) \leftarrow \lambda(q_j)$ 
             $\lambda_e(q_{j2}) \leftarrow \lambda(q_j)$ 
        else
           $\delta_e(q_i, s) \leftarrow q_j$  for every  $s \in t_j^i$ 
           $\lambda_e(q_j) \leftarrow \lambda(q_j)$ 


```

Algorithm 3: Extended LTL₃ monitor Construction

Automata-based Synchronous Monitoring

Example

Consider the LTL₃ monitor for $\varphi = \mathbf{F}(a \wedge b)$.



The set of outgoing transitions from monitor state q_0 is $T_0 = \{t_0^0, t_+^0\}$ where:


$$t_0^0 = \{\{a\}, \{b\}, \emptyset\}$$

$$t_+^0 = \{\{a, b\}\}$$

We can verify that t_0^0 is indistinguishable from t_+^0 . Therefore we split transition t_0^0 into two transitions $t_{01}^0 = \{\{a\}\}$ and $t_{02}^0 = \{\{b\}, \emptyset\}$.



Example




Note that we have

$$\lambda(q_{01}) = \lambda(q_0) = ?$$

$$\lambda(q_{02}) = \lambda(q_0) = ?$$

Example




Note that we have

$$\lambda(q_{01}) = \lambda(q_0) = ?$$

$$\lambda(q_{02}) = \lambda(q_0) = ?$$

Example



Note that we have

$$\lambda(q_{01}) = \lambda(q_0) = ?$$

$$\lambda(q_{02}) = \lambda(q_0) = ?$$

Automata-based Synchronous Monitoring

Example

Suppose $s = \{a, b\}$ is the current global state of the system and Let $\varphi = \mathbf{F}(a \wedge b)$. The initial state of the monitors is as follows:

sample				round 1	
	a	b	LS_i^1		LS_i^2
M_1	<i>true</i>	\perp	$\{q_{02}, q_{\top}\}$	M_1	crashed
M_2	\perp	<i>true</i>	$\{q_{01}, q_{\top}\}$	M_2	$\{q_{\top}\}$
M_3	\perp	\perp	$\{q_{01}, q_{02}, q_{\top}\}$	M_3	$\{q_{01}, q_{\top}\}$
M_4	\perp	\perp	$\{q_{01}, q_{02}, q_{\top}\}$	M_4	$\{q_{01}, q_{\top}\}$

round 2		round 3	
	LS_i^3		LS_i^4
M_1	crashed	M_1	crashed
M_2	crashed	M_2	crashed
M_3	$\{q_{\top}\}$	M_3	$\{q_{\top}\}$
M_4	$\{q_{01}, q_{\top}\}$	M_4	$\{q_{\top}\}$

Presentation outline

- 1 Motivation
- 2 LTL₃ Monitor
- 3 Problem Statement
- 4 Related Work
- 5 Contributions
- 6 Synchronous Monitoring
- 7 Asynchronous Monitoring**
- 8 Conclusion
- 9 Future Work

Asynchronous Monitoring

Asynchronous Monitoring

The system under inspection produces a finite trace $\alpha = s_0 s_1 \cdots s_k$, and is inspected with respect to an LTL formula φ by a set $\mathcal{M} = \{M_1, M_2, \dots, M_n\}$ of asynchronous distributed monitors.

Local Monitor Algorithm

Data: Extended LTL₃ monitor $\mathcal{M}_e^\varphi = \{\Sigma, Q_e, q_0, \delta_e, \lambda_e\}$ and state s_j


Result: a set of monitor states

- 1 Let $\mathcal{S}_i^{s_j}$ be the initial concrete local state of the monitor ;
- 2 $Snap_i^{s_j} \leftarrow \mathcal{S}_i^{s_j}$
- 3 emit $V_i^{s_j} = \mu_2(\mu_1(Snap_i^{s_j}, \mathcal{M}^\varphi))$

Asynchronous Monitoring

Example

Let $\varphi = \mathbf{F}(a \wedge b)$ whose Extended LTL₃ monitor is given below. Suppose monitors are at monitor state q_0 , and let $s = \{a, b\}$ be the global state of the system.





Example

The following tables represent each monitor M_i 's initial local snapshot $Snap_i^s$ and its verdict set V_i calculated based on only $Snap_i^s$.

$$Snap_1^s$$

	M_1	M_2	M_3
a	<i>true</i>	\perp	\perp
b	\perp	\perp	\perp
V_1	$\{q_1, q_\top\}$		

$$Snap_2^s$$

	M_1	M_2	M_3
a	\perp	\perp	\perp
b	\perp	<i>true</i>	\perp
V_2	$\{q_0, q_\top\}$		

$$Snap_3^s$$

	M_1	M_2	M_3
a	\perp	\perp	\perp
b	\perp	\perp	\perp
V_3	$\{q_0, q_1, q_\top\}$		

$$V_1 \cap V_2 \cap V_3 = q_\top.$$

Presentation outline

- 1 Motivation
- 2 LTL₃ Monitor
- 3 Problem Statement
- 4 Related Work
- 5 Contributions
- 6 Synchronous Monitoring
- 7 Asynchronous Monitoring
- 8 Conclusion**
- 9 Future Work

Conclusion

Conclusion

- We proposed a synchronous monitoring algorithm that copes with f crash failures in a distributed setting. The algorithm solves the synchronous monitoring problem in $f + 1$ rounds of communication and reduces the message size overhead from $|AP|$ to $\log(m_q)$.
- We proposed an algorithm for distributed crash-resilient asynchronous RV that **consistently** monitors the system under inspection with **no communication** between monitors.

Presentation outline

- 1 Motivation
- 2 LTL₃ Monitor
- 3 Problem Statement
- 4 Related Work
- 5 Contributions
- 6 Synchronous Monitoring
- 7 Asynchronous Monitoring
- 8 Conclusion
- 9 Future Work

Future Work

Futur Work

- To address more severe faults, e.g., Byzantine failures.
- To have monitors observe, communicate, and emit verdicts between any two global states.
- To extend our results to the case where the input to the monitors is a sequence of global states and each monitor produces a sequence of verdict sets, one per each global state